

# FIT2086 Lecture 8

## Model Selection and Penalized Regression

Daniel F. Schmidt

Faculty of Information Technology, Monash University

September 10, 2018

## 1 Model Selection Criteria

- Overfitting and Underfitting
- Model Selection Methods

## 2 Penalized Regression Techniques

- Statistical Stability
- Penalized Regression

# Revision from last week (1)

- We also have  $p$  predictor variables  $X_1, \dots, X_p$
- We can build a classifier for  $Y$  using our predictors, i.e., we want to find

$$\mathbb{P}(Y = y \mid X_1 = x_1, X_2 = x_2, \dots, X_p = x_p)$$

- Naïve Bayes assumes predictors are conditionally independent, given the value of the target

$$P(x_1, \dots, x_p \mid y) = \prod_{j=1}^p P(x_j \mid y)$$

- Naïve Bayes classifier:

$$P(y \mid x_1, \dots, x_p) = \frac{\prod_{j=1}^p P(x_j \mid y) P(y)}{P(x_1, \dots, x_p)}$$

# Revision from last week (2)

- A logistic regression models the conditional log-odds as

$$\log \left( \frac{\mathbb{P}(Y_i = 1 \mid x_{i,1}, \dots, x_{i,p})}{\mathbb{P}(Y_i = 0 \mid x_{i,1}, \dots, x_{i,p})} \right) = \beta_0 + \sum_{j=1}^p \beta_j x_{i,j} \equiv \eta_i$$

- Logistic regression model of conditional probability

$$\mathbb{P}(Y_i = 1 \mid x_{i,1}, \dots, x_{i,p}) = \frac{1}{1 + \exp(-\eta_i)}$$

- Estimated using maximum likelihood
- Performance measures for classification
  - Classification error
  - Sensitivity and specificity
  - Area-under-the-curve (AUC)
  - Logarithmic loss

## 1 Model Selection Criteria

- Overfitting and Underfitting
- Model Selection Methods

## 2 Penalized Regression Techniques

- Statistical Stability
- Penalized Regression

# Supervised Learning – recap

- Imagine we have measured  $p + 1$  variables on  $n$  individuals (people, objects, things)
- One variable is our target
- We would like to predict our target using the remaining  $p$  variables (predictors)
- If the variable we are predicting is categorical, we are performing **classification**
- If the variable we are predicting is numerical, we are performing **regression**

# Underfitting/Overfitting Example (1)

- We often have many measured predictors
  - Should we use them all, and if not, why not?
- Omitting important predictors
  - Called **underfitting**
  - Leads to systematic error (“bias”) in predicting the target
- Including spurious predictors
  - Called **overfitting**
  - Leads our model to “learn” noise and random variation
  - Poorer ability to predict to new, unseen data from our population

# Underfitting/Overfitting Example (2)

- Example: we observe  $x$  and  $y$  data and want to build a prediction model for  $y$  using  $x$ 
  - Data looks nonlinear so we use polynomial regression
  - We take  $x, x^2, x^3, \dots, x^{20} \Rightarrow$  very flexible model
  - How many terms to include?
- For example, do we use

$$y = \beta_0 + \beta_1 x + \beta_2 x^2 + \varepsilon$$

or


$$y = \beta_0 + \beta_1 x + \beta_2 x^2 + \beta_3 x^3 + \beta_4 x^4 + \beta_5 x^5 + \varepsilon$$

or another model with some other number of polynomial terms.




# Underfitting/Overfitting Example (3)

- Example dataset of 50 samples




# Underfitting/Overfitting Example (4)

- Use  $(x, x^2)$ , too simple – underfitting




# Underfitting/Overfitting Example (5)

- Use  $(x, x^2, \dots, x^{20})$ , too complex – overfitting



# Underfitting/Overfitting Example (6)

- $(x, x^2, \dots, x^5)$  seems “just right”. But how to find this model?



# Mean-Squared Prediction Error (1)

- Introduce the idea of a **true**, underlying (population) model  
 $\implies$  this is the model that “generated” the data
- For regression:

$$\mathbb{E}[Y \mid x_1, \dots, x_p] = f(x_1, \dots, x_p)$$

- For binary classification:

$$\log \left( \frac{\mathbb{P}(Y = 1 \mid x_1, \dots, x_p)}{1 - \mathbb{P}(Y = 1 \mid x_1, \dots, x_p)} \right) = f(x_1, \dots, x_p)$$

- In real life,  $f(\cdot)$  could be any function
- For simplicity, let's focus on regression

# Mean-Squared Prediction Error (2)

- We receive a sample of data from our population  $\mathbf{y} = (y_1, \dots, y_n)$ 
  - This is our **training data**
- We use this to build a model  $\hat{\mathcal{M}}(\mathbf{y}) \equiv \hat{\mathcal{M}}$  to estimate  $f(\cdot)$
- In linear regression, a model  $\mathcal{M}$  is a particular linear combination of
  - predictors  $x_1, \dots, x_p$ , and
  - non-linear transformations of the predictors
- We would like our model to predict well onto new data
  - We call this the model's ability to **generalise**
- How can we measure this?

# Mean-Squared Prediction Error (3)

- Imagine we receive new data  $\mathbf{y}' = (y'_1, \dots, y'_m)$  from our population
  - This is called **testing data**
- Once we have a model,  $\hat{\mathcal{M}}(\mathbf{y})$ , we can make predictions

$$\hat{y}_i(\hat{\mathcal{M}}(\mathbf{y})) \equiv \hat{y}(x'_{i,1}, \dots, x'_{i,p}, \hat{\mathcal{M}}(\mathbf{y}))$$

for the value of  $y'_i$ , given predictors  $x'_{i,1}, \dots, x'_{i,p}$ :

- We can measure how well our model predicts the new data using

$$\text{MSPE}(\hat{\mathcal{M}}(\mathbf{y})) = \frac{1}{m} \sum_{i=1}^m (y'_i - \hat{y}_i(\hat{\mathcal{M}}(\mathbf{y})))^2$$

which is the mean-squared prediction error on our testing data

# Mean-Squared Prediction Error (4)

- Remember that our predictions  $\hat{y}(x_1, \dots, x_p, \hat{\mathcal{M}}(\mathbf{y}))$  are produced by a model that we learned using the training data  $\mathbf{y}$   
 $\Rightarrow$  they depend on the training data that we have observed
- If we sampled **different** training data from our population, we would learn a different model,  $\hat{\mathcal{M}}(\mathbf{y})$ , and get different predictions  $\hat{y}_i(\hat{\mathcal{M}}(\mathbf{y}))$
- To remove this dependency, we can *average* the  $\text{MSPE}(\hat{\mathcal{M}}(\mathbf{y}))$  over all the different models we could learn from all the different possible training samples we could draw from our population
- We can write this *expected* mean-squared prediction error as

$$\text{EMSPE} = \mathbb{E} [\text{MSPE}(\hat{\mathcal{M}}(\mathbf{y}))] = \text{bias}^2 + \text{variance}$$

where

- bias measures **systematic error** of our predictions;
- variance measures the error incurred through **random variation**



# Mean-Squared Prediction Error (4)

- Remember that our predictions  $\hat{y}(x_1, \dots, x_p, \hat{\mathcal{M}}(\mathbf{y}))$  are produced by a model that we learned using the training data  $\mathbf{y}$   
 $\Rightarrow$  they depend on the training data that we have observed
- If we sampled **different** training data from our population, we would learn a different model,  $\hat{\mathcal{M}}(\mathbf{y})$ , and get different predictions  $\hat{y}_i(\hat{\mathcal{M}}(\mathbf{y}))$
- To remove this dependency, we can *average* the  $\text{MSPE}(\hat{\mathcal{M}}(\mathbf{y}))$  over all the different models we could learn from all the different possible training samples we could draw from our population
- We can write this *expected* mean-squared prediction error as

$$\text{EMSPE} = \mathbb{E} [\text{MSPE}(\hat{\mathcal{M}}(\mathbf{y}))] = \text{bias}^2 + \text{variance}$$

where

- bias measures **systematic error** of our predictions;
- variance measures the error incurred through **random variation**

# Underfitting/Overfitting (1)

- We can demonstrate these ideas on our example
- The true model for this example was

$$\mathbb{E}[y | x] = f(x) = 9.7x^5 + 0.8x^3 + 9.4x^2 - 5.7x - 2$$

- To generate a sample from this true model:
  - I random sampled  $\mathbf{x} = (x_1, \dots, x_{50})$  uniformly from  $(-1, 1)$
  - For each  $x_i$  value, I generated the samples  $y_i$  using

$$y_i = f(x_i) + \varepsilon$$


where  $\varepsilon_i \sim N(0, 1)$ .

# Underfitting and Bias (1)

- When **underfitting**, our model is too simple to capture the truth
- In linear regression, this means we are omitting important predictors
- For our example, imagine we fit a quadratic polynomial to our data
  - The “truth” is a fifth-order polynomial
- Regardless of how much training data from the population we see, we cannot make a quadratic look like a fifth order polynomial
- The inability of our model to fit the truth well leads to systematic error in predicting  $f(\cdot)$
- The increase in EMSPE introduced by underfitting is therefore primarily due to **bias**


# Underfitting and Bias (2)

- Using  $\mathcal{M} = \{x, x^2\}$  on a new sample – too simple




# Underfitting and Bias (3)

- Again, using  $\mathcal{M} = \{x, x^2\}$  – again, too simple



# Underfitting and Bias (4)

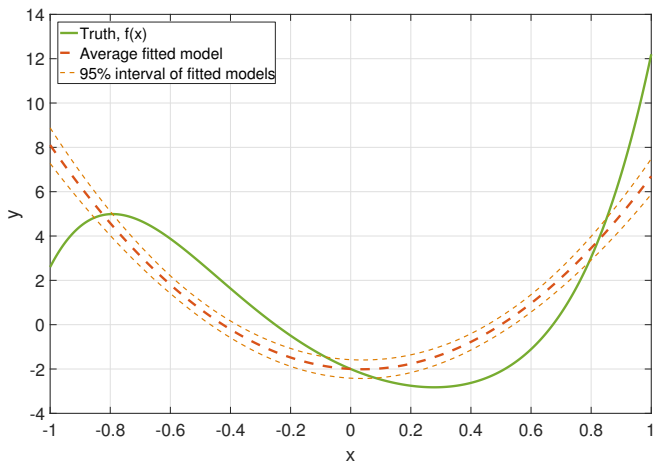
- Again, using  $\mathcal{M} = \{x, x^2\}$  – once more, too simple



# Underfitting and Bias (5)

- Bias and variance were determined by **simulation**
- Generated 10,000 samples of size  $n = 50$  from our true model (our training data)
- For each sample, fitted a 2nd order polynomial, i.e.,  $x, x^2$  terms
- Produced predictions for 5,000 equally spaced  $x'$  values in  $(-1, 1)$  (our testing data) for each of the 10,000 different models
  - Calculated the average prediction at each  $x'$  value
  - Calculated the variance of predictions at each  $x'$  value
  - Used these to get  $\text{bias}^2$  and variance
- For visualisation:
  - Calculated the average prediction at each  $x'$  value
  - Calculated 2.75 and 97.5 percentiles of predictions at each  $x'$

# Underfitting and Bias (6)



Fitting 2nd order polynomial to data generated from a 5th order polynomial; average fit is poor (high bias) but does not vary greatly (low variance).

$\text{EMSPE} = 3.33$ ,  $\text{bias}^2 = 3.27$ ,  $\text{variance} = 0.058$




# Overfitting and Variance (1)

- When **overfitting**, our model is more complex than the truth
- In linear regression, this means we are including **unassociated** predictors
- The extra predictor introduces a relationship between predictor and target that is not really there  
⇒ we learn random “noise” from our training data
- In our example, imagine we fitted a 7th order polynomial to our data
  - The “truth” is a 5th order polynomial, so the coefficients for the 6th and 7th order polynomial terms should be zero
  - But when we learn them from our training data, due to random chance, they will never be exactly equal to zero
- The increase in EMSPE introduced by overfitting is therefore primarily due to **variance**


# Overfitting and Variance (2)

- Using  $(x, x^2, \dots, x^7)$  on a new sample – too complex




# Overfitting and Variance (3)

- Again, using  $(x, x^2, \dots, x^7)$  – too complex




# Overfitting and Variance (4)

- Again, using  $(x, x^2, \dots, x^7)$  – too complex




# Overfitting and Variance (5)



Fitting 7th order polynomial to data generated from a 5th order polynomial; average fit is good (no bias) but does varies more than 2nd order polynomial.

EMSPE = 0.1470,  $\text{bias}^2 = 0$ , variance = 0.1470

# Overfitting and Variance (6)



Fitting 20th order polynomial to data generated from a 5th order polynomial; average fit is good (no bias) but variation is higher, especially near boundaries.

EMSPE = 0.4860,  $\text{bias}^2 = 0$ , variance = 0.4860

# Underfitting/Overfitting (1)

- In general:
  - Simpler models have higher bias, lower variance
  - More complex models have lower bias, higher variance
  - Variance decreases with increasing  $n$ , bias does not
  - Variance increases with increasing number of predictors
- Recall EMSPE can be written as:

$$\text{EMSPE} = \text{bias}^2 + \text{variance}$$

- Model selection for prediction is about trading these off  
⇒ select the right level of complexity for the data
- In linear/logistic regression, “complexity” is increased by including more predictors, or transformations of predictors
  - Omitting a predictor will reduce variance, but may lead to bias if predictor is important

# Underfitting/Overfitting (2)

- Model selection is quite subtle
- In linear models, including all predictors that are associated with target may not be enough to remove bias
  - If relationship between target and predictor is nonlinear, also need to include the right transformations ...
- You can overfit and underfit at the same time
  - Omit important predictor, include unassociated predictor ...
- Sometimes can do better by omitting associated predictors!
  - If association is weak and noise is large, may be too hard to get a good estimate of the coefficient  $\beta_j$
  - Ergo, may be better to omit it
- Combining tools with real world knowledge/thinking can often improve predictions



# Hypothesis Testing to Select Predictors (1)

- Recap: can select predictors using hypothesis testing  
 $\implies$  we know that a predictor  $j$  is unimportant if  $\beta_j = 0$
- So we can test the hypothesis:

$$H_0 \quad : \quad \beta_j = 0$$

vs

$$H_A \quad : \quad \beta_j \neq 0$$

which, in this setting is a variant of the  $t$ -test (see Ross, Chapter 9 and Studio 6)

- Imagine we tried fitting each predictor by itself to the data
  - Could choose to include predictor if  $p$ -value was less than 0.05

# Hypothesis Testing to Select Predictors (2)

- If we were testing one hypothesis, this would guarantee probability of false discovery to be 5%
  - If predictor is unassociated, we would incorrectly conclude it **to be associated** for 5% of possible samples from our population
- But if we have  $p$  predictors, we need to do  $p$  tests
- So, even if all predictors were unassociated, would expect  $0.05p$  tests to have  $p$ -values  $< 0.05$  just by chance
- If  $p$  is large, this can be big, i.e., if  $p = 1000$ , we would falsely believe 50 predictors were associated just by chance on average
- This can easily happen if we include lots of transformations; e.g., trying all squares, logarithms, interactions of predictors
- Sometimes called **data dredging**

# Hypothesis Testing to Select Predictors (3)

- This is called the **multiple testing problem**
- We examine one way to resolve it: the **Bonferroni procedure**
- Let  $\alpha$  be our significance level
- This says we should instead reject null hypothesis only if

$$p\text{-value} < \frac{\alpha}{p}$$

- So if  $\alpha = 0.05$  and  $p = 1000$ , we would need to see  $p$ -values of  $0.05/1000 = 5 \times 10^{-5}$  to believe that a predictor is associated with the target

# Hypothesis Testing to Select Predictors (4)

- Bonferroni guarantees the probability of making **at least one** false discovery is 5%
- Very conservative approach, but can be confident that predictors we identify are very likely truly associated
- Related method called **false discovery rate** (FDR)  
⇒ more power to discover associated predictors

# Likelihood

- The likelihood  $p(\mathbf{y} | \hat{\mathcal{M}})$  is a general measure of goodness of fit
- The likelihood quantifies how much probability our model assigns to the data we have observed  
 $\implies$  more probability, the more compatible the data is with our model
- Given a probability distribution  $p(y | \theta)$  and estimates  $\hat{\theta}$  for the parameters, compute likelihood by:
  - 1 Plug our estimated parameters  $\hat{\theta}$  into our probability distribution
  - 2 Evaluate the probability assigned to each  $y_1, \dots, y_n$  by this model, i.e.,

$$p(\mathbf{y} | \hat{\theta}) = \prod_{i=1}^n p(y_i | \hat{\theta})$$

- We have seen likelihoods of
  - Simple Gaussian, Poisson, Bernoulli models (Lecture 3, Studio 3)
  - Linear regression models (Lecture 6)
  - Logistic regression models (Lecture 7)

# Model Selection Criteria (1)

- We can use adjusted likelihood to guide selection of models  
⇒ can control complexity through model selection criteria
- A model selection criterion assigns a score to a model which
  - takes into account how well the model fits the data (likelihood)
  - takes into account how complex the model is
- Search for the model that minimises model selection criterion  
⇒ in linear/logistic regression, find good set of predictors
- Let
  - $\mathcal{M}$  denote a model (set of predictors to use);
  - $\hat{\mathcal{M}}$  denote the model fitted to data  $\mathbf{y}$ ;
  - $L(\mathbf{y} | \hat{\mathcal{M}})$  denote the minimised negative log-likelihood for the model  $\mathcal{M}$
- Neg-log-likelihood **always decreases** with more predictors

## Model Selection Criteria (2)

- Penalize neg-log-likelihood of model by a complexity penalty
- The **Akaike information criterion (AIC)**:

$$\text{AIC}(\hat{\mathcal{M}}) = L(\mathbf{y} | \hat{\mathcal{M}}) + k_{\mathcal{M}}$$

where  $k_{\mathcal{M}}$  is number of predictors in model  $\mathcal{M}$

- The **Kullback–Leibler information criterion (KIC)**:

$$\text{KIC}(\hat{\mathcal{M}}) = L(\mathbf{y} | \hat{\mathcal{M}}) + 3/2 k_{\mathcal{M}}$$

- Both of these tend to overfit, even for large  $n$ 
  - Moderate overfitting less damaging for large  $n$ ;
  - Both rarely underfit
- In practice, KIC often performs better than AIC

## Model Selection Criteria (2)

- Penalize neg-log-likelihood of model by a complexity penalty
- The **Akaike information criterion (AIC)**:

$$\text{AIC}(\hat{\mathcal{M}}) = L(\mathbf{y} | \hat{\mathcal{M}}) + k_{\mathcal{M}}$$

where  $k_{\mathcal{M}}$  is number of predictors in model  $\mathcal{M}$

- The **Kullback–Leibler information criterion (KIC)**:

$$\text{KIC}(\hat{\mathcal{M}}) = L(\mathbf{y} | \hat{\mathcal{M}}) + 3/2 k_{\mathcal{M}}$$

- Both of these tend to overfit, even for large  $n$ 
  - Moderate overfitting less damaging for large  $n$ ;
  - Both rarely underfit
- In practice, KIC often performs better than AIC



# Model Selection Criteria (3)

- The **Bayesian information criterion (BIC)**:

$$\text{BIC}(\hat{\mathcal{M}}) = L(\mathbf{y} | \hat{\mathcal{M}}) + k_{\mathcal{M}}/2 \log n$$

where  $n$  is the number of samples used to fit  $\hat{\mathcal{M}}$

- As  $n$  increases, probability of overfitting goes to zero
  - More conservative; for small  $n$  can sometimes underfit
- The **Risk inflation information criterion (RIC)**:

$$\text{RIC}(\hat{\mathcal{M}}) = L(\mathbf{y} | \hat{\mathcal{M}}) + k_{\mathcal{M}} \log p$$

where  $p$  is total number of possible predictors

- Similar to multiple testing idea; can be quite conservative
  - Recommended over AIC/KIC if  $p$  is very large
- Many refined criteria exist: MML, MDL, corrected AIC/KIC  
 $\Rightarrow$  often better for small  $n$

# Model Selection Criteria (3)

- The **Bayesian information criterion (BIC)**:

$$\text{BIC}(\hat{\mathcal{M}}) = L(\mathbf{y} | \hat{\mathcal{M}}) + k_{\mathcal{M}}/2 \log n$$

where  $n$  is the number of samples used to fit  $\hat{\mathcal{M}}$

- As  $n$  increases, probability of overfitting goes to zero
  - More conservative; for small  $n$  can sometimes underfit
- The **Risk inflation information criterion (RIC)**:

$$\text{RIC}(\hat{\mathcal{M}}) = L(\mathbf{y} | \hat{\mathcal{M}}) + k_{\mathcal{M}} \log p$$

where  $p$  is total number of possible predictors

- Similar to multiple testing idea; can be quite conservative
  - Recommended over AIC/KIC if  $p$  is very large
- Many refined criteria exist: MML, MDL, corrected AIC/KIC  
 $\Rightarrow$  often better for small  $n$

# Cross validation (1)

- **Cross-validation (CV)** is a very general strategy
- Motivation: ideally, we would select model that minimises prediction error on future data  $y'_1, \dots, y'_m$

$$\text{MSPE}(\hat{\mathcal{M}}(\mathbf{y})) = \frac{1}{m} \sum_{i=1}^m (y'_i - \hat{y}_i(\hat{\mathcal{M}}(\mathbf{y})))^2$$

⇒ requires knowing our testing data ahead of time ...

- We don't have future testing data, but we know the data  $\mathbf{y}$  we used to fit the model is generated by the same process
- So we can try and **estimate** the MSPE


# Cross validation (2)

- To estimate MSPE using cross validation
  - 1 We randomly partition data  $\mathbf{y}$  we have into two sets:
    - A training set  $\mathbf{y}_{\text{train}}$ ,
    - and a testing set  $\mathbf{y}_{\text{test}}$
  - 2 Fit a model  $\mathcal{M}$  to the training data  $\mathbf{y}_{\text{train}}$
  - 3 Compute the MSPE of this model on the testing data  $\mathbf{y}_{\text{test}}$
  - 4 Repeat this procedure multiple times and average the MSPE
- This gives us an estimate of how well our model would predict future data from the same population
- Do this for all our different models, and choose the one that has the smallest CV error

# Cross validation (3)

- $K$ -fold CV
  - split into  $K$  equal sized random partitions;
  - train on  $K - 1$  partitions, test on the remaining partition
  - repeat this  $m$  times
- The bigger  $m$ , the more stable the estimate
- **Leave-one-out CV (LOO CV):**
  - train on  $n - 1$  of the samples, test on the remaining sample
  - do this for all  $n$  possible partitions
- Cross-validation is:
  - simple to implement, very flexible;
  - potential slow if model fitting is computationally expensive
- LOO CV is similar to AIC for large samples  $n$   
 $\Rightarrow$  can be very different for small samples

# Example: fitting a polynomial



Model selection scores for fitting zero-order through twentieth-order polynomials to data generated by our example fifth-order polynomial. The likelihood always decreases with extra polynomial terms. The dots signify the models selected by each criterion.

## 1 Model Selection Criteria

- Overfitting and Underfitting
- Model Selection Methods

## 2 Penalized Regression Techniques

- Statistical Stability
- Penalized Regression

# Conventional Model Selection Methods (1)

- **All-subsets selection:**
  - Try all combination of predictors to model with smallest model selection criterion score
- **Forward selection algorithm:**
  - 1 Start with the empty model;
  - 2 Find the predictor that reduces info criterion by most
  - 3 If no predictor improves model, end.
  - 4 Add this predictor to the model
  - 5 Return to Step 2
- **Backwards selection** is related algorithm
  - Start with the full model and remove predictors



# Conventional Model Selection Methods (2)

- Problems with traditional methods
- **Statistically unstable**
  - Small changes in data  $\Rightarrow$  big changes in model
- Multiple hypothesis testing problem
  - Potential issues with false positives
- All-subsets selection
  - Infeasible for moderate to large number of predictors  $p$
- Stepwise selection
  - Affected adversely by **correlation in predictors**
  - Slow for large  $p$

# Statistical Instability (1)

- Example dataset ( $n = 354$ ,  $p = 10$ )
  - Target:  $Y$ , measure of diabetes progression
  - Predictors:
    - AGE, SEX, BMI
    - BP: blood pressure
    - $S_1, \dots, S_6$ : blood serum measurements
- Experiment
  - ① Removed one of the 354 samples at random
  - ② Found “best” combination of predictors
    - Checked all possible (1024) combinations of predictors
    - Selected combination with smallest CV
- Repeated this four times

## Statistical Instability (2)

Test	AGE	SEX	BMI	BP	S1	S2	S3	S4	S5	S6
1		✓	✓	✓	✓	✓	✓		✓	
2		✓	✓	✓	✓	✓	✓			
3	✓	✓	✓	✓	✓	✓	✓	✓		
4	✓	✓	✓	✓	✓	✓	✓		✓	

Predictors chosen by selecting model that minimises cross-validation score using all-subsets selection method. For each test, a single one of the  $n = 354$  samples was removed at random.

- The “all-or-nothing” nature of this type of model selection (either exclude or include predictor) means that the “best” model can change substantially with only a small change to the data.

# Penalized Regression (1)

- Alternative: apply a “penalty” to the coefficients
- For linear regression, we have penalized least-squares:

$$\left(\hat{\beta}_0, \hat{\beta}_\lambda\right) = \arg \min_{\beta_0, \beta} \left\{ \text{RSS}(\beta_0, \beta) + \lambda \sum_{j=1}^p g(\beta_j) \right\}$$

- In comparison regular least squares only minimises goodness-of-fit (RSS)
- How to choose the penalty function  $g(\cdot)$ ?
  - Chosen to increase with increasing magnitude of  $\beta_j$
- Generalises easily to logistic regression, i.e., penalized maximum likelihood

# Penalized Regression (2)

- Intuition:
  - Large coefficients imply strong effects
  - Coefficient of zero implies no association
- For this to make sense we need to **standardise** our predictors before fitting, i.e., ensure that

$$\sum_{i=1}^n x_{i,j} = 0 \text{ and } \sum_{i=1}^n x_{i,j}^2 = n$$

- This can be easily done by subtracting mean and dividing by standard deviation
- Each predictor is then on the same (arbitrary) scale  
 $\Rightarrow$  implies that larger  $\beta_j$  means stronger association

# Penalized Regression (3)

- The penalty includes a user-chosen **hyperparameter**  $\lambda$
- $\lambda$  controls how *strong* the penalty is
- Generally,
  - As  $\lambda$  goes to zero, the estimates of  $\beta$  are closer to least-squares
  - As  $\lambda$  gets larger, the estimates of  $\beta$  become smaller
- Choosing  $\lambda$  sets the “complexity” of the regression model
- By varying  $\lambda$  we can generate a “path” of regression models
  - The models become more complex as we move along the path

# Penalized Regression (4)

- Advantages of penalized regression:
  - Increased statistical stability
  - Can be applied even when the sample size is smaller than the number of predictors (the  $n < p$  setting)
  - Handles correlation between predictors
- **Multi-collinearity** occurs when at least two of the predictors are highly correlated
  - This means it is difficult to distinguish between their relationship with the target
  - In traditional least-squares this leads to increased variance
  - In stepwise methods, this dramatically increases instability

# Ridge Regression (1)

- The first penalized method introduced
- **Ridge regression** solves


$$\left(\hat{\beta}_0, \hat{\beta}_\lambda\right) = \arg \min_{\beta_0, \beta} \left\{ \text{RSS}(\beta_0, \beta) + \lambda \sum_{j=1}^p \beta_j^2 \right\}$$

- Strengths:
  - Very quick to run, even for large  $n$  and  $p$
  - Produces very stable, low variance estimates
  - Very appropriate for multi-collinearity
- Weaknesses:
  - Cannot estimate coefficients to be exactly zero
    - i.e., no variable selection




## Ridge Regression (2)

- Penalty is the square of the coefficient ( $\ell_2$  penalty)



# Ridge Regression (3)

- Example ridge regression coefficient path



# Lasso Regression (1)


- A “sparse” estimator
- Lasso regression solves

$$\left(\hat{\beta}_0, \hat{\beta}_\lambda\right) = \arg \min_{\beta_0, \beta} \left\{ \text{RSS}(\beta_0, \beta) + \lambda \sum_{j=1}^p |\beta_j| \right\}$$

- Strengths:
  - Efficient algorithms exist
  - Can estimate coefficients to be exactly zero
    - i.e., can perform variable selection
- Weaknesses:
  - Produces biased estimates for large coefficients
  - Correlated predictors can be problematic
  - Can overfit (include unassociated predictors)


# Lasso Regression (2)

- Penalty is the absolute value of the coefficient ( $\ell_1$  penalty)



# Lasso Regression (3)

- Example lasso regression coefficient path




# Selecting a Model Using Penalized Regression

- How to choose the hyperparameter  $\lambda$
- Standard procedure:
  - 1 Vary  $\lambda$  over some grid of values
  - 2 For each  $\lambda$ , use cross-validation to estimate prediction error
  - 3 Select  $\lambda$  with smallest cross-validation error
  - 4 Use that  $\lambda$  to estimate our final model from all the data
- Can also use methods like AIC, BIC, etc.
- Implementations:
  - MATLAB `lasso()` and `lassoglm()` functions
  - R, use the `glmnet` package


# Example – Lasso regression

- Example: we observe  $x$  and  $y$  data and want to build a prediction model for  $y$  using  $x$ 
  - Data looks nonlinear so we use polynomial regression
  - We take  $x, x^2, x^3, \dots, x^{20} \Rightarrow$  very flexible model
  - Which terms to include?
- For example, our model could be:
  - $(x, x^2, x^3)$ ; or
  - $(x, x^4, x^{16}, x^{17})$ ; or
  - $(x^2, x^3, x^5, x^6, x^9)$ ; and so on ...
- Let us use Lasso regression

# Example – Lasso regression




(a) Fitted model




(b) Lasso path



# Example – Lasso regression




(a) Fitted model




(b) Lasso path

# Example – Lasso regression




(a) Fitted model




(b) Lasso path

# Example – Lasso regression




(a) Fitted model




(b) Lasso path

# Example – Lasso regression




(a) Fitted model




(b) Lasso path

# Example – Lasso regression




(a) Fitted model




(b) Lasso path

# Example – Lasso regression




(a) Fitted model




(b) Lasso path

# Example – Lasso regression




(a) Fitted model




(b) Lasso path

# Example – Lasso regression




(a) Fitted model




(b) Lasso path



# Example – Lasso regression



(a) Fitted model



(b) Lasso path

# Bias/Variance Trade-off

- Penalized regression has a bias/variance motivation
- Recall that the EMSPE of a model can be written as

$$\text{EMSPE} = \text{bias}^2 + \text{variance}$$

- If all relevant predictors included, least-squares is unbiased  
⇒ however, LS variance can be high
- Penalized regression introduce some bias but reduces variance
  - They **shrink** the coefficients towards zero
- Choosing a good  $\lambda$  means we can reduce the overall EMSPE

- Terms you should know:
  - Underfitting and bias
  - Overfitting and variance
  - Expected mean-squared prediction error
  - Multiple hypothesis testing, Bonferroni procedure
  - AIC, BIC
  - Cross-validation
  - Statistical instability
  - Penalized regression
  - Ridge/lasso regression
- Next week: non-linear machine learning algorithms for classification and regression ( $k$  nearest neighbours, decision trees, random forests).