

# FIT2086 Lecture 10

## Introduction to Unsupervised Learning

Daniel F. Schmidt

Faculty of Information Technology, Monash University

October 1, 2018

## 1 Clustering/Mixture Modelling

- Clustering
- Mixture Modelling

## 2 Matrix Completion

- Matrix Completion Problem
- Methods for Matrix Completion

# Revision from last week (1)

- Machine learning methods
- Cross validation for model selection
  - Withhold data to estimate prediction error
  - $K$ -fold CV divides data up into  $K$  equal sized groups
  - Train on  $K - 1$  folds, predict on the remaining fold
- Decision Trees
  - Split the data up by asking questions of the predictors
  - Number of leaves determines complexity of tree
  - Easy to interpret, flexible
- Methods for learning trees
  - Greedy growing of trees – find best split at each step
  - Backwards pruning of large tree
  - Use CV to select number of leaves in the tree

# Revision from last week (2)

- Trees have low bias, high variance
- One solution: random forests
  - Grow many trees with guided random search
  - Aggregate predictions from the trees
  - Stable, low variance, but loses interpretability
- $k$  nearest neighbours (kNN) methods
  - Assume individuals similar in predictors are similar in targets
  - Find  $k$  “most similar” individuals in data to new individual
  - Use their targets to predict target for new individual
- Use CV to select  $k$ , other tuning parameters

# Assignment 2 (1)

- Bonus question asked to make predictions
- Use concrete data to learn a model of compressive strength of concrete, and use it to predict compressive strength of new concrete mixes
- 28 people submitted predictions
- Methods people tried:
  - Interactions
  - Non-linear transformations of predictors (logs, polynomials)
  - Pruning using the R `step()` function
  - Pruning using change in  $R^2$  value
  - Combining with non-linear models such as trees

## Assignment 2 (2)

### Concrete Prediction Results

Who	Score	Notes
24 people	$> 126.36$	Various methods
Me	126.36	Basic step-wise linear model
Me	124.17	Full linear model (least squares)

# Assignment 2 (2)

## Concrete Prediction Results

Who	Score	Notes
24 people	$> 126.36$	Various methods, mostly overfitting
Me	126.36	Basic step-wise linear model
Me	124.17	Full linear model (least squares)
Chang Yu Wang	109.40	Combination of $k$ -NN and random forest

# Assignment 2 (2)

## Concrete Prediction Results

Who	Score	Notes
24 people	$> 126.36$	Various methods, mostly overfitting
Me	126.36	Basic step-wise linear model
Me	124.17	Full linear model (least squares)
Chang Yu Wang	109.40	Combination of $k$ -NN and random forest
Andy Ying Zhi Chong	105.33	Additive polynomial model with order selection



# Assignment 2 (2)

## Concrete Prediction Results

Who	Score	Notes
24 people	$> 126.36$	Various methods, mostly overfitting
Me	126.36	Basic step-wise linear model
Me	124.17	Full linear model (least squares)
Chang Yu Wang	109.40	Combination of $k$ -NN and random forest
Andy Ying Zhi Chong	105.33	Additive polynomial model with order selection
Niky Chen	103.47	Square of Age and binarization of two variables

# Assignment 2 (2)

## Concrete Prediction Results

Who	Score	Notes
24 people	$> 126.36$	Various methods, mostly overfitting
Me	126.36	Basic step-wise linear model
Me	124.17	Full linear model (least squares)
Chang Yu Wang	109.40	Combination of $k$ -NN and random forest
Andy Ying Zhi Chong	105.33	Additive polynomial model with order selection
Niky Chen	103.47	Square of Age and binarization of two variables
Me	86.60	Additive model with splines

# Assignment 2 (2)

## Concrete Prediction Results

Who	Score	Notes
24 people	$> 126.36$	Various methods, mostly overfitting
Me	126.36	Basic step-wise linear model
Me	124.17	Full linear model (least squares)
Chang Yu Wang	109.40	Combination of $k$ -NN and random forest
Andy Ying Zhi Chong	105.33	Additive polynomial model with order selection
Niky Chen	103.47	Square of Age and binarization of two variables
Me	86.60	Additive model with splines
Ryan Whitelock-Jones	80.80	Combined random forest with carefully constructed linear model of transformed variables

# Assignment 2 (2)

## Concrete Prediction Results

Who	Score	Notes
24 people	> 126.36	Various methods, mostly overfitting
Me	126.36	Basic step-wise linear model
Me	124.17	Full linear model (least squares)
Chang Yu Wang	109.40	Combination of $k$ -NN and random forest
Andy Ying Zhi Chong	105.33	Additive polynomial model with order selection
Niky Chen	103.47	Square of Age and binarization of two variables
Me	86.60	Additive model with splines
Ryan Whitelock-Jones	80.80	Combined random forest with carefully constructed linear model of transformed variables
Me (cheating)	72.89	Full linear model + Age <sup>1/3</sup>

## 1 Clustering/Mixture Modelling

- Clustering
- Mixture Modelling

## 2 Matrix Completion

- Matrix Completion Problem
- Methods for Matrix Completion

# Unsupervised learning (1)

- We have  $n$  items, each with  $q$  associated attributes, formed into a matrix

$$\mathbf{Y} = \begin{pmatrix} \mathbf{y}_1 \\ \mathbf{y}_2 \\ \vdots \\ \mathbf{y}_n \end{pmatrix} = \begin{pmatrix} y_{1,1} & y_{1,2} & \cdots & y_{1,q} \\ y_{2,1} & y_{2,2} & \cdots & y_{2,q} \\ \vdots & \vdots & \ddots & \vdots \\ y_{n,1} & y_{n,2} & \cdots & y_{n,q} \end{pmatrix}$$

- Each  $\mathbf{y}_i$  is a “data-point” in  $q$ -dimensional space
- Unlike supervised learning, we do not nominate any one of these as a “target”
- Instead we want to discover structure in the data

# Unsupervised learning (2)

- What is unsupervised learning used for?
- Classifying or categorising objects (taxonomy)
  - For example, species of animals
- Filling in missing entries in the data matrix
  - Matrix completion problem
  - Recommender systems
  - Imputation (estimating missing data in predictor matrix before supervised learning)
- Image processing
  - Noise removal
  - Compression
  - Image analysis and recognition

# Unsupervised learning vs supervised learning

- Supervised learning: target  $Y$  and explanatory variables  $X_1, \dots, X_p$ 
  - We then try and find the conditional distribution

$$p(Y | X_1, \dots, X_p)$$

using a specific form of model (linear regression, tree, etc.)

- Model describes relationship between  $Y$  and  $X_1, \dots, X_p$
- Unsupervised learning: only have explanatory variables  $X_1, \dots, X_q$ 
  - We try and discover the joint distribution

$$p(X_1, \dots, X_q)$$

using a specific form of model

- The details of the model reveal internal structure of data



- Assumptions
  - Population consists of  $K$  sub-populations ( $K > 1$ )
  - We are given observations from the pooled population only
    - No sub-population information is available
- Aim
  - Discover the number of sub-populations  $K$
  - Estimate models for each of the sub-populations
- Sometimes called **intrinsic classification**  
⇒ Class labels are learned from the data

# K-means Clustering (1)

- Perhaps most commonly used clustering technique
- Models data as having  $K$  “centroids” defined by mean vectors

$$\mathbf{M} = \begin{pmatrix} \boldsymbol{\mu}_1 \\ \vdots \\ \boldsymbol{\mu}_K \end{pmatrix} = \begin{pmatrix} \mu_{1,1} & \cdots & \mu_{1,q} \\ \vdots & \ddots & \vdots \\ \mu_{K,1} & \cdots & \mu_{K,q} \end{pmatrix}$$

- Assigns items to class with most similar mean vector
- Similarity between item  $i$  and centroid  $k$  is

$$d_k(i) = \left( \sum_{j=1}^q (y_{i,j} - \mu_{k,j})^2 \right)^{\frac{1}{2}}$$

⇒ Euclidean distance between the vectors.

# $K$ -means Clustering (1)

- Perhaps most commonly used clustering technique
- Models data as having  $K$  “centroids” defined by mean vectors

$$\mathbf{M} = \begin{pmatrix} \boldsymbol{\mu}_1 \\ \vdots \\ \boldsymbol{\mu}_K \end{pmatrix} = \begin{pmatrix} \mu_{1,1} & \cdots & \mu_{1,q} \\ \vdots & \ddots & \vdots \\ \mu_{K,1} & \cdots & \mu_{K,q} \end{pmatrix}$$


- Assigns items to class with most similar mean vector
- Similarity between item  $i$  and centroid  $k$  is

$$d_k(i) = \left( \sum_{j=1}^q (y_{i,j} - \mu_{k,j})^2 \right)^{\frac{1}{2}}$$

$\Rightarrow$  Euclidean distance between the vectors.

# K-means Clustering (2)


- Artificial data example



- Chosen so that the “clusters” are obvious for demonstration purposes

# K-means Clustering (3)

- K-means clustering with  $K = 2$




- Centroids chosen to minimise the within-cluster sum-of-squares

# $K$ -means Algorithm (1)

- The  $k$ -means algorithm is very simple:
  - ① Initialise  $\mu_1, \dots, \mu_K$  randomly
  - ② Loop until convergence
    - ① Compute distances  $d_k(i)$  from each data point  $\mathbf{y}_i$  to each centroid  $\mu_k$
    - ② Assign datapoints to cluster with closest centroid
    - ③ Re-estimate each  $\mu_k$  using the datapoints assigned to cluster  $k$
- Converges quickly to a stable solution  
⇒ might not be the global-minima
- Sensitive to starting points


# K-means Algorithm (2)

- Example:  $K = 3$ , initial starting points for centroids  $\mu_k$



# K-means Algorithm (3)


- Example: assigning points to clusters with closest centroid






# K-means Algorithm (4)

- Example: re-estimating centroids from data in the clusters




# K-means Algorithm (5)

- Example: assigning points to clusters with closest centroid




# K-means Algorithm (6)

- Example: re-estimating centroids from data in the clusters




# K-means Algorithm (7)

- Example: after 3 iterations, centroids are stable



# $K$ -means Algorithm (9)

- The  $k$ -means algorithm is sensitive to starting points



# K-means Algorithm (10)

- $k$ -means tries to optimise the function

$$D(\boldsymbol{\mu}_1, \dots, \boldsymbol{\mu}_K) = \sum_{i=1}^n \min_k \{d_k(i)\}$$

$\implies$  Tries to minimise distance of each point to nearest centroid

- Bad seeding leads to local minima
- $k$ -means++ algorithm improves convergence dramatically
  - Randomly choose centers to be far apart from each other

# Further Clustering

- Alternative similarity measures
  - Weighted Euclidean distance
  - “Cityblock” distance
  - Hamming distance (for pure binary data)
  - and many more ...
- Some potential issues
  - “Hard” classification of items to clusters
  - Difficult to handle mixed attributes (continuous, discrete)
  - No explicit statistical interpretation
  - How to choose  $K$  using just the data?
- Mixture modelling a flexible alternative

# Further Clustering

- Alternative similarity measures
  - Weighted Euclidean distance
  - “Cityblock” distance
  - Hamming distance (for pure binary data)
  - and many more ...
- Some potential issues
  - “Hard” classification of items to clusters
  - Difficult to handle mixed attributes (continuous, discrete)
  - No explicit statistical interpretation
  - How to choose  $K$  using just the data?
- Mixture modelling a flexible alternative



# Mixture Modelling (1)

- Models data as a **mixture of probability distributions**

$$p(y_{i,j}) = \sum_{k=1}^K \alpha_k p(y_{i,j} \mid \theta_{k,j})$$

where

- $K$  is the number of classes
  - $\alpha = (\alpha_1, \dots, \alpha_K)$  are the mixing (population) weights
  - $\theta_{k,j}$  are the parameters of the distributions
- Has an explicit probabilistic form  
 $\implies$  allows for statistical interpretation

# Mixture Modelling (2)


- How is this related to clustering?
- Each class is a cluster
  - Class-specific probability distributions over each attribute
    - e.g., normal, inverse Gaussian, Poisson, etc.
  - Mixing weight is prevalence of items in the class
    - Fraction of our population in that particular subpopulation
- The resulting mixture model has
  - $K$  different classes (subpopulations)
  - $q$  different models for each class, one for each attribute
    - $\theta_{k,j}$  are parameters of model for attribute  $j$  in class  $k$
  - $K \times q$  total probability models

# Mixture Modelling (2)

- How is this related to clustering?
- Each class is a cluster
  - Class-specific probability distributions over each attribute
    - e.g., normal, inverse Gaussian, Poisson, etc.
  - Mixing weight is prevalence of items in the class
    - Fraction of our population in that particular subpopulation
- The resulting mixture model has
  - $K$  different classes (subpopulations)
  - $q$  different models for each class, one for each attribute
    - $\theta_{k,j}$  are parameters of model for attribute  $j$  in class  $k$
  - $K \times q$  total probability models

# Mixture Modelling (3)

- Example: two normal distributions



# Mixture Modelling (4)

- Measure of similarity of item to class

$$p_k(\mathbf{y}_i) = \prod_{j=1}^q p(y_{i,j} \mid \boldsymbol{\theta}_{k,j})$$

⇒ probability of item's attributes under class distributions

- For Gaussian models, this is equivalent to Euclidean distance
- For non-Gaussian models (Bernoulli, Poisson, etc.) it is often a generalisation of the Euclidean distance
  - Related to something called Kullback–Leibler divergence

# Mixture Modelling (5)

- Membership of items to classes is **soft**

$$r_{i,k} = \frac{\alpha_k p_k(\mathbf{y}_i)}{\sum_{l=1}^K \alpha_l p_l(\mathbf{y}_i)}$$

- Application of Bayes' theorem
- Posterior probability of belonging to class  $k$ 
  - $\alpha_k$  is a *a priori* probability item belongs to class  $k$
  - $p_k(\mathbf{y}_i)$  is probability of data item  $\mathbf{y}_i$  under class  $k$ $\implies$  Assign to class with highest posterior probability
- Total number of samples in a class is then

$$n_k = \sum_{i=1}^n r_{i,k}$$

# Mixture Modelling (5)

- Membership of items to classes is **soft**

$$r_{i,k} = \frac{\alpha_k p_k(\mathbf{y}_i)}{\sum_{l=1}^K \alpha_l p_l(\mathbf{y}_i)}$$

- Application of Bayes' theorem
- Posterior probability of belonging to class  $k$ 
  - $\alpha_k$  is a *a priori* probability item belongs to class  $k$
  - $p_k(\mathbf{y}_i)$  is probability of data item  $\mathbf{y}_i$  under class  $k$ $\implies$  Assign to class with highest posterior probability
- Total number of samples in a class is then

$$n_k = \sum_{i=1}^n r_{i,k}$$

# Multivariate Normal Distribution (1)

- So far we have considered separate univariate distributions for each attribute
- However, it would be useful to model attributes as related
- **Multivariate normal distributions** are important in statistics
  - Generalize normal distributions to more than one dimension
  - Allow for correlation between random variables
- Are important in mixture model
- They model relationships between multiple random variables
  - The attributes of an individual are likely related
  - For example, height and weight will show correlation



# Multivariate Normal Distribution (2)

- If  $\mathbf{Y} = (Y_1, \dots, Y_q)$  are RVs with pdf

$$\left(\frac{1}{2\pi}\right)^{\frac{q}{2}} \sqrt{|\Sigma^{-1}|} \exp\left(-\frac{1}{2}(\mathbf{y} - \boldsymbol{\mu})' \Sigma^{-1}(\mathbf{y} - \boldsymbol{\mu})\right)$$

then they are multivariate normal with means  $\boldsymbol{\mu} = (\mu_1, \dots, \mu_q)$  and covariance matrix  $\Sigma$


- The entry  $\mu_j$  is the mean for  $Y_j$
- The entry

$$\Sigma_{i,j} = \text{cov}(Y_i, Y_j)$$

is the covariance between  $Y_i$  and  $Y_j$ .


# Multivariate Normal Distribution (3)

- Example,  $\boldsymbol{\mu} = (0, 0)$ ,  $\boldsymbol{\Sigma} = \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix}$




# Multivariate Normal Distribution (4)

- Example,  $\boldsymbol{\mu} = (2, 0)$ ,  $\boldsymbol{\Sigma} = \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix}$




# Multivariate Normal Distribution (5)

- Example,  $\boldsymbol{\mu} = (0, 0)$ ,  $\boldsymbol{\Sigma} = \begin{pmatrix} 2 & 0 \\ 0 & 1 \end{pmatrix}$




# Multivariate Normal Distribution (6)

- Example,  $\boldsymbol{\mu} = (0, 0)$ ,  $\boldsymbol{\Sigma} = \begin{pmatrix} 1 & 0 \\ 0 & 1.5 \end{pmatrix}$




# Multivariate Normal Distribution (7)

- Example,  $\boldsymbol{\mu} = (0, 0)$ ,  $\boldsymbol{\Sigma} = \begin{pmatrix} 1 & 0.6 \\ 0.6 & 1 \end{pmatrix}$



# Multivariate Normal Distribution (8)

- Example,  $\boldsymbol{\mu} = (0, 0)$ ,  $\boldsymbol{\Sigma} = \begin{pmatrix} 1 & -0.9 \\ -0.9 & 1 \end{pmatrix}$



# Multivariate Normal Distribution (9)

- Multivariate normal generalises the univariate normal distribution
  - For  $q = 1$ , reduces to usual normal distribution
- Several different common covariance structures:
  - Diagonal  $\Sigma$ , all variances the same (spherical)
  - Diagonal  $\Sigma$ , variances differing
  - Arbitrary  $\Sigma$  (elliptical)
- Each structure has more parameters to estimate  
 $\implies$  more flexible, but more complex



# Estimating Mixture Models (1)

- Given class memberships, the negative log-likelihood of data in class  $k$  is

$$-\sum_{i=1}^n r_{i,k} \sum_{j=1}^q \log p(y_{i,j} | \theta_{k,j})$$

$\implies$  **weighted** negative log-likelihood

- Use **expectation-maximisation** (EM) algorithm
  - Estimate parameters,  $\theta_{k,j}$ , ( $k = 1, \dots, K$ ), ( $j = 1, \dots, q$ ) using weighted maximum likelihood
  - Re-calculate class memberships  $r_{i,k}$  based on new parameters
  - If estimates have not stabilised, go to step (1)
- Initialise model with random class memberships
- Generalisation of  $k$ -means

# Estimating Mixture Models (1)

- Given class memberships, the negative log-likelihood of data in class  $k$  is

$$-\sum_{i=1}^n r_{i,k} \sum_{j=1}^q \log p(y_{i,j} | \theta_{k,j})$$

$\Rightarrow$  **weighted** negative log-likelihood

- Use **expectation-maximisation** (EM) algorithm
  - Estimate parameters,  $\theta_{k,j}$ , ( $k = 1, \dots, K$ ), ( $j = 1, \dots, q$ ) using weighted maximum likelihood
  - Re-calculate class memberships  $r_{i,k}$  based on new parameters
  - If estimates have not stabilised, go to step (1)
- Initialise model with random class memberships
- Generalisation of  $k$ -means

# Estimating Mixture Models (2)

- Find  $K$  by minimising a goodness-of-fit criterion
- Difficult, non-convex optimisation problem  
⇒ **Many local minima**
- Each iteration, do the following
  - Remove classes with too few data points
  - Attempt to split all classes
  - Attempt to combine pairs of classes
  - Randomly assign data to classes, and re-estimate
- The mixture model with the smallest criterion score is retained, and the process is repeated

# Estimating Mixture Models (2)

- Find  $K$  by minimising a goodness-of-fit criterion
- Difficult, non-convex optimisation problem  
⇒ Many local minima
- Each iteration, do the following
  - Remove classes with too few data points
  - Attempt to split all classes
  - Attempt to combine pairs of classes
  - Randomly assign data to classes, and re-estimate
- The mixture model with the smallest criterion score is retained, and the process is repeated

# Estimating Mixture Models (3)

- **Information Criteria** goodness-of-fit criterion
  - Popular for learning mixture models
- **Information criterion** score is our yardstick, comprised of
  - 1 Goodness of fit of the mixture model to the data
  - 2 Model complexity penalty based on number of classes/parameters

⇒ choose model which balances complexity against fit
- Popular method is called minimum message length
  - Developed here at Monash by C.S.Wallace
  - Uses information theory interpretation of probability
  - Compress data using model; find model that leads to shortest compressed data


# Estimating Mixture Models (3)

- **Information Criteria** goodness-of-fit criterion
  - Popular for learning mixture models
- **Information criterion** score is our yardstick, comprised of
  - 1 Goodness of fit of the mixture model to the data
  - 2 Model complexity penalty based on number of classes/parameters

⇒ choose model which balances complexity against fit
- Popular method is called minimum message length
  - Developed here at Monash by C.S.Wallace
  - Uses information theory interpretation of probability
  - Compress data using model; find model that leads to shortest compressed data


# Example (1)

- Example: two dimensional dataset



## Example (2)


- Mixture modelling discovers  $K = 4$  classes






## Example (3)

- Plot of the mixture model density




# Example: Galaxy data (1)



Data on  $n = 82$  galaxies; each data point is the velocity of a galaxy.


## Example: Galaxy data (2)



Mixture modelling finds  $K = 4$  classes.

8.9%  $N(9.71, 0.2)$ , 23%  $N(19.74, 0.3)$ , 62%  $N(22, 5.25)$ , 4%  $N(33.04, 1.27)$

## Example: Galaxy data (3)



Mixture modelling finds  $K = 4$  classes.

8.9%  $N(9.71, 0.2)$ , 23%  $N(19.74, 0.3)$ , 62%  $N(22, 5.25)$ , 4%  $N(33.04, 1.27)$

# Example: Multivariate Data Analysis (1)

- Well known diabetes dataset
  - 268 diabetics, 500 non-diabetics
  - 768 samples, with 8 predictors
  - 763 missing exposure measurements (12%)
- Outcome is diabetes in Pima indians (DIA)

Pima Indians Variables

	Name	Mean	$\sigma$	Min	Max	% Missing
	Number of Pregnancies (PREG)	4.5	3.2	1	17	14.4%
	Plasma Glucose Concentration (PLAS)	121.6	30.5	44	199	0.6%
	Diastolic Blood Pressure (BP)	72.4	12.4	24	122	4.5%
	Triceps Skin Fold Thickness (SKIN)	29.1	10.5	7	99	29.5%
	2-hour Serum Insulin (INS)	155.5	118.8	14	846	48.7%
	Body Mass Index (BMI)	32.4	6.9	18.2	67.1	1.4%
	Diabetes Pedigree Function (PED)	0.47	0.33	0.078	2.42	0%
	Age (AGE)	33.2	11.7	21	81	0%

# Example: Multivariate Data Analysis (2)

- Estimate mixture model for exposures and outcome
  - All predictors Gaussian, target (diabetes) is Bernoulli
  - $I_4 = 18,719.1$ ,  $I_5 = 18,713.0$ ,  $I_6 = 18,714.7$ ,  $I_7 = 18,732.7$

Pima Indians Mixture Model (Means)

Class	$\hat{\alpha}_k$	PREG	PLAS	BP	SKIN	INS	BMI	PED	AGE	DIA
1	0.13	2.5	150	75	35	238	37	0.59	33	0.82
2	0.23	7.6	141	78	33	214	35	0.52	43	0.78
3	0.25	2.0	104	66	20	105	27	0.42	24	0.02
4	0.19	2.7	112	71	34	138	36	0.47	26	0.20
5	0.18	6.4	110	75	28	117	30	0.41	42	0.06

## 1 Clustering/Mixture Modelling

- Clustering
- Mixture Modelling

## 2 Matrix Completion

- Matrix Completion Problem
- Methods for Matrix Completion

# Matrix Completion Problem (1)

- We have a large matrix of data  $\mathbf{Y}$ 
  - Rows of  $\mathbf{Y}$  are individuals
  - Columns of  $\mathbf{Y}$  are attributes of individuals
- Many entries of  $\mathbf{Y}$  are missing
  - Usually they are unmeasured
- **Matrix completion** involves filling in the missing entries
- Assume individuals are independent, attributes are dependent
  - Use dependencies between attributes to estimate missing entries



# Matrix Completion Problem (2)

- Some applications of matrix completion

- ① **Imputation**

- Matrix of features for a supervised learning problem
    - Most supervised learning methods cannot handle missing data
    - Filling in missing entries lets us use entire matrix

- ② **Recommender systems**

- Matrix is set of ratings/purchasers
    - Rows are individuals, columns are products
    - For example, Netflix challenge

	Terminator	Love Actually	Aliens	Predator	Bridesmaids
	2	4	2	1	5
	4	1	5	4	1
	4	—	4	—	—
	—	4	—	—	—

- Estimate missing ratings and recommend movies

# Matrix Completion Problem (3)

- Univariate imputation

- Simplest approach to imputation
- Estimate a statistical model each column
- Replace missing entries with suitable statistic
  - Mean/median for numeric variables
  - Mode for categorical variables
- Ignores structure and relationships between variables
- Is very fast

- Multivariate normal

- Specify correlations between variables
- Estimate missing entries using correlation info
- Takes into account relationships between variables
- Assumes data is clustered in one single cluster
  - Can use mixture modelling to extend this idea further

# Imputation using $k$ -Nearest Neighbours (1)

- Recall  $k$ -nearest neighbours algorithm
  - We have a set of  $n$  example predictor/target pairs
    - Predictor values  $x_{i,1}, \dots, x_{i,p}$  paired with target  $y_i$
  - We want to predict target value for new individual with predictor values  $x'_1, \dots, x'_p$
- Find  $k$  individuals in our data “most similar” to the new individual
  - Use target values of these  $k$  individuals to predict target for our new individual
- Very weak assumptions
  - Individuals similar to each other in terms of predictor values will be similar in terms of targets
- Use cross-validation to select neighbourhood size  $k$

# Imputation using $k$ -Nearest Neighbours (2)

- Easily adapted to matrix completion
- When computing similarity, ignore missing values
- Then, for each column  $j = 1, \dots, p$ 
  - Predict each missing entry in column  $j$  using all other columns as explanatory variables
- Sometimes called **collaborative filtering**
- Netflix example using  $k = 1$

Terminator	Love Actually	Aliens	Predator	Bridesmaids
2	4	2	1	5
4	1	5	4	1
4	?	4	—	—
—	4	—	—	—

# Imputation using $k$ -Nearest Neighbours (2)

- Easily adapted to matrix completion
- When computing similarity, ignore missing values
- Then, for each column  $j = 1, \dots, p$ 
  - Predict each missing entry in column  $j$  using all other columns as explanatory variables
- Sometimes called **collaborative filtering**
- Netflix example using  $k = 1$

Terminator	Love Actually	Aliens	Predator	Bridesmaids
2	4	2	1	5
4	1	5	4	1
4	?	4	—	—
—	4	—	—	—

# Imputation using $k$ -Nearest Neighbours (3)

- Easily adapted to matrix completion
- When computing similarity, ignore missing values
- Then, for each column  $j = 1, \dots, p$ 
  - Predict each missing entry in column  $j$  using all other columns as explanatory variables
- Sometimes called **collaborative filtering**
- Netflix example using  $k = 1$

Terminator	Love Actually	Aliens	Predator	Bridesmaids
2	4	2	1	5
4	1	5	4	1
4	?	4	—	—
—	4	—	—	—

# Imputation using $k$ -Nearest Neighbours (4)

- Easily adapted to matrix completion
- When computing similarity, ignore missing values
- Then, for each column  $j = 1, \dots, p$ 
  - Predict each missing entry in column  $j$  using all other columns as explanatory variables
- Sometimes called **collaborative filtering**
- Netflix example using  $k = 1$

Terminator	Love Actually	Aliens	Predator	Bridesmaids
2	4	2	1	5
4	1	5	4	1
4	1	4	—	—
—	4	—	—	—

- Terms you should know:
  - Clustering
  - $k$ -means algorithm
  - Mixture modelling
  - Matrix completion
  - Imputation
  - Collaborative filtering
- Next week: simulation based methods (bootstrapping, permutation tests, random number generation)